

Executing and Optimizing Hadoop MapReduce

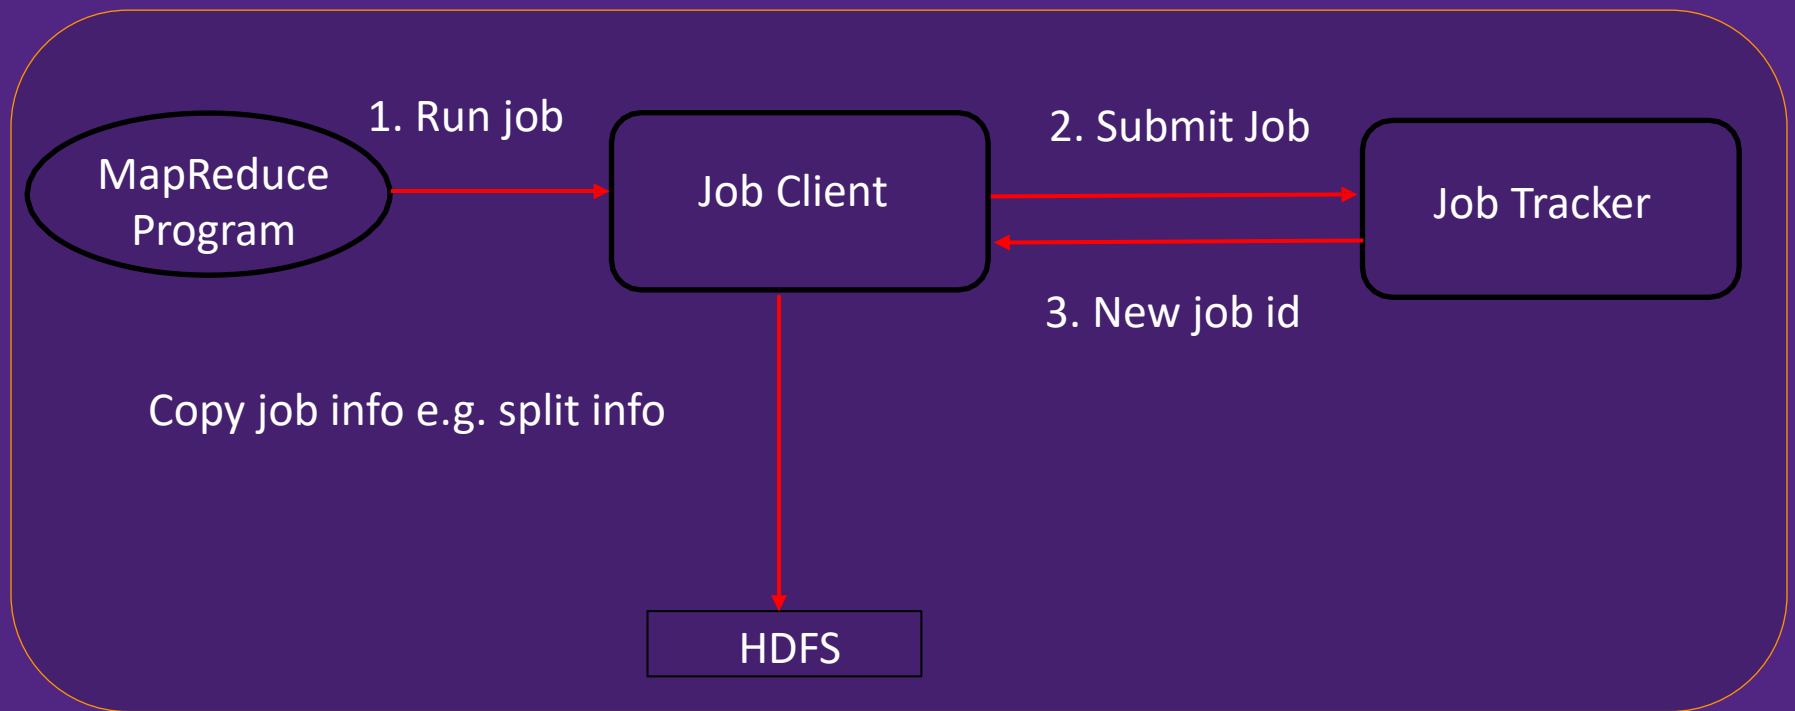
CS 4417B

The University of Western Ontario

Terminology

- A **job** is a “full program” with mapper and reducer code, input dataset and a location for the output
- A **task** is an execution of a mapper or a reducer on part of the data
- Example: Run “Word Count” across a file with 20 blocks is one **job**
 - Assume a file is assigned to one map **task**
 - 20 blocks typically result in 20 map **tasks**
 - These may send data to many reduce **tasks**

Execution Architecture



- **JobClient**
 - Submit job request

Execution Architecture

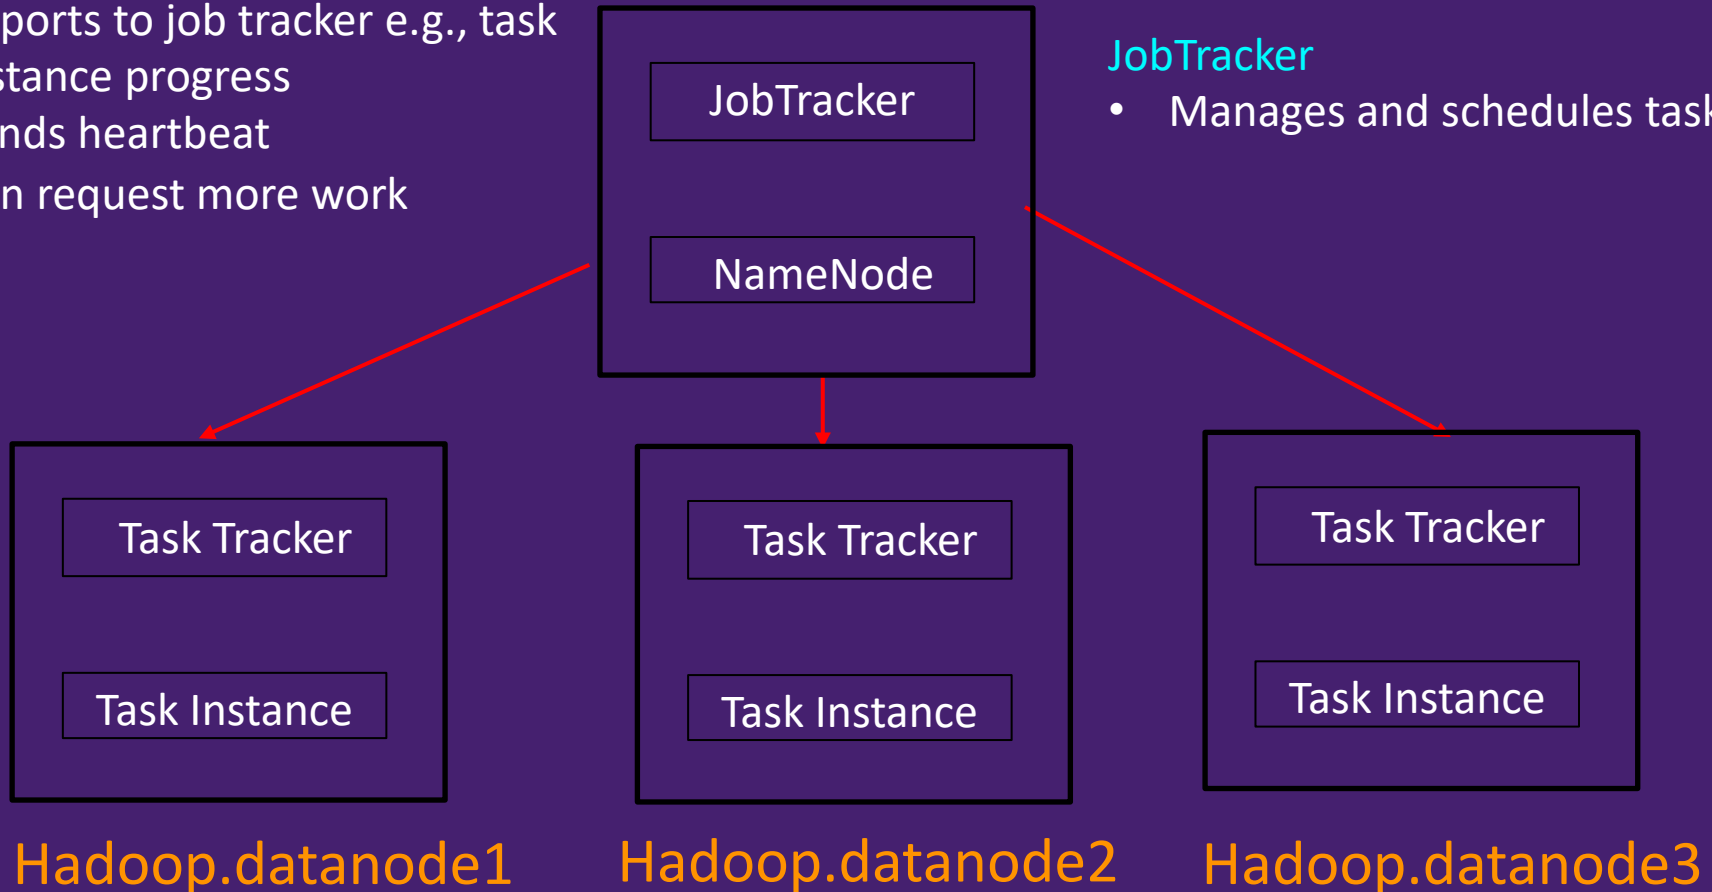
Hadoop.namenode

Task Tracker

- Reports to job tracker e.g., task instance progress
- Sends heartbeat
- Can request more work

JobTracker

- Manages and schedules tasks



Job Processing

- **JobTracker** assigns work to TaskTrackers
- After map, the **TaskTrackers** exchange map-output information to build the reduce key space
- **JobTracker** partitions **reduce()** keyspace into **m** chunks, where **m** is set by user
 - Assigns work to reducers

Assigning to Reducers

- Given a set of keys $\{k_1 \dots k_n\}$ how do we assign?
- Use a $\text{hash}(k_i)$
 - Actually $\text{hash}(k_i) \bmod R$ where R is number of reducers
- Issues?
 - Assumes equal distribution of data over hashes (that's why hashes work well)

How many Mappers and Reducers

- Mappers
 - By default this is the number of **HDFS blocks** being processed
 - The number of maps can also be controlled by specifying the **minimum split size**
- Reducers
 - This is a function of the number of nodes and the amount of data – for **n** nodes; two suggestions:
 - $0.95n$ (all reduces start immediately)
 - $1.75n$ (fast nodes can end up doing 2 reduce jobs)

References

- Data-Intensive Text Processing with MapReduce
 - Jimmy Lin and Chris Dyer
 - <http://lintool.github.io/MapReduceAlgorithms/MapReduce-book-final.pdf>
- Hadoop docs:
<https://hadoop.apache.org/docs/stable/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html>

Optimizing using Local Aggregation: "Combiners"

Optimization

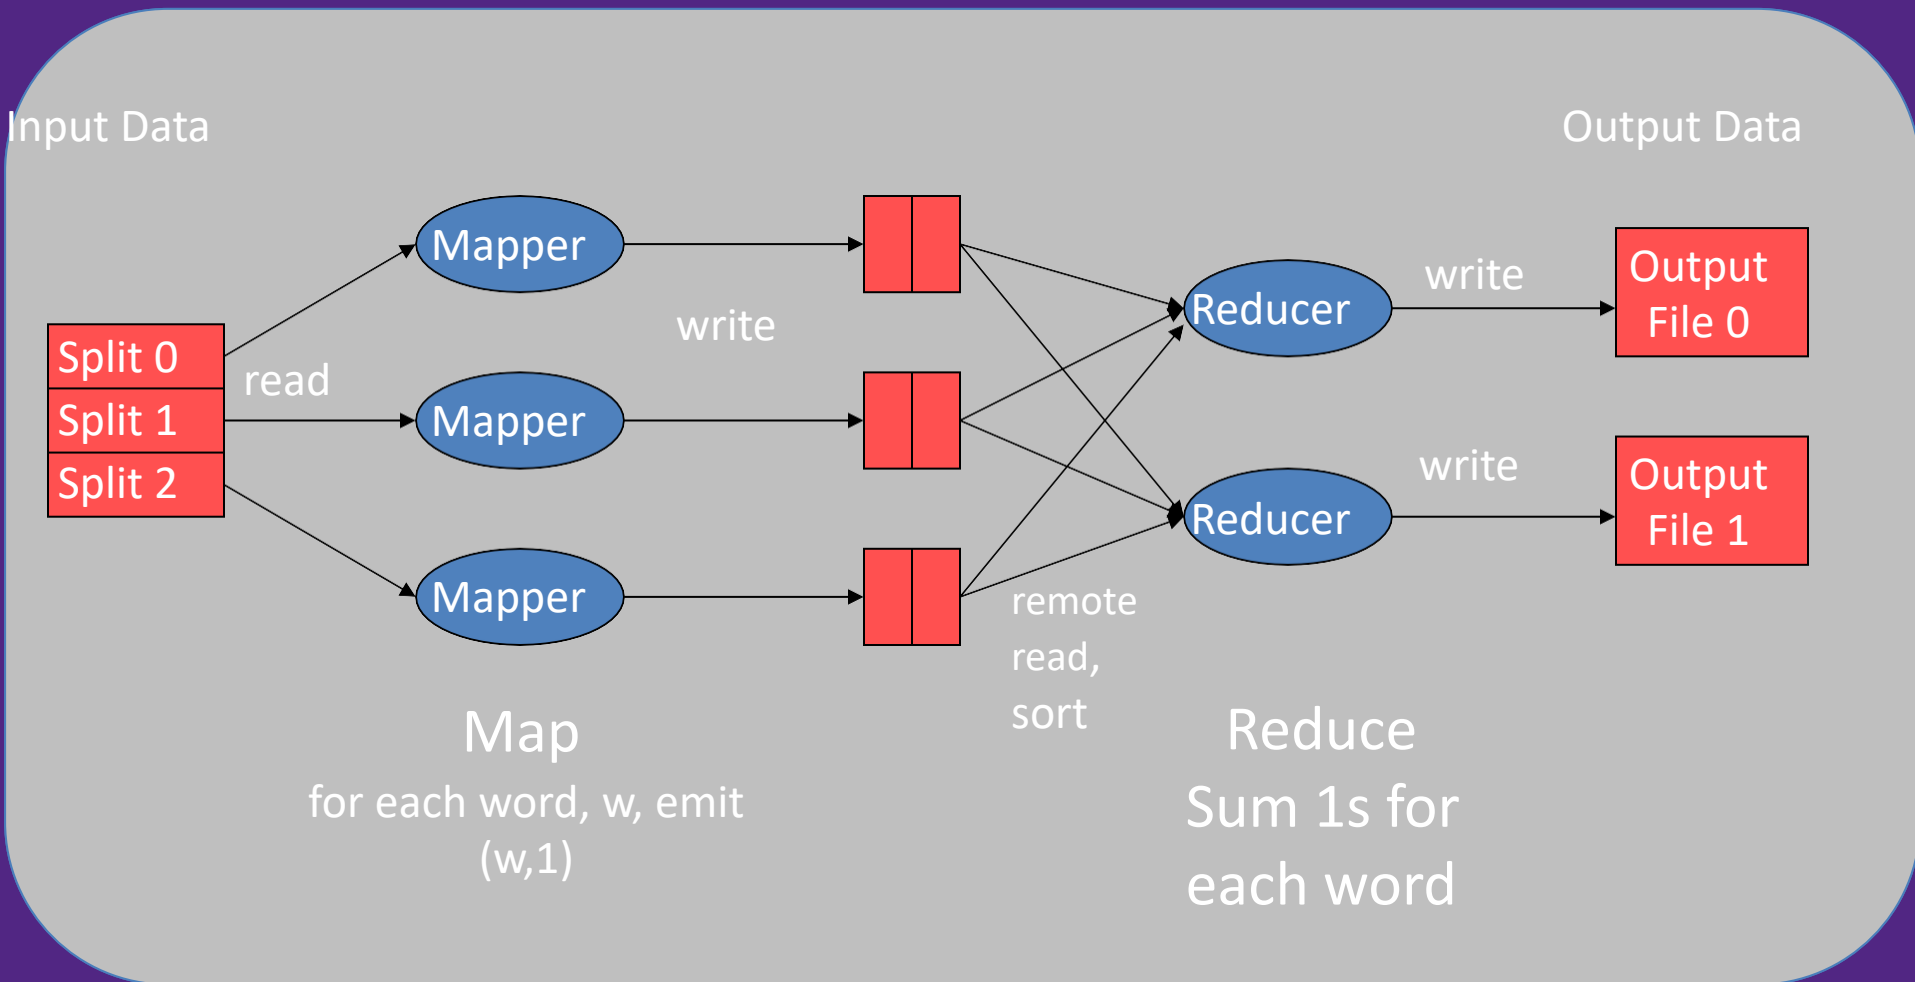
"Plain" MapReduce

- Map
 - One input pair -> one output pair
- Reduce
 - Many input pairs with same key -> one output pair
- By loosening this definition, we can improve performance

Revisiting Word Count

- Map function
 - For each word, **w**, generate (**w**,1)
- Reduce function
 - Reducers sum the ones for each word
 - Output: For each word, **w**,
emit (**w**, number of times w occurs)

MapReduce Workflow



Revisiting Word Count

- Issue: Large number of key-value pairs
 - Example: (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1)
(w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1)
(w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1)
(w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1) (w,1)
(w,1) (w,1) (w,1) (w,1)
- The output of the mapper is written to disk and read from the disk by reducer
- Writing and reading to disks is relatively slow
- Potential for network congestion

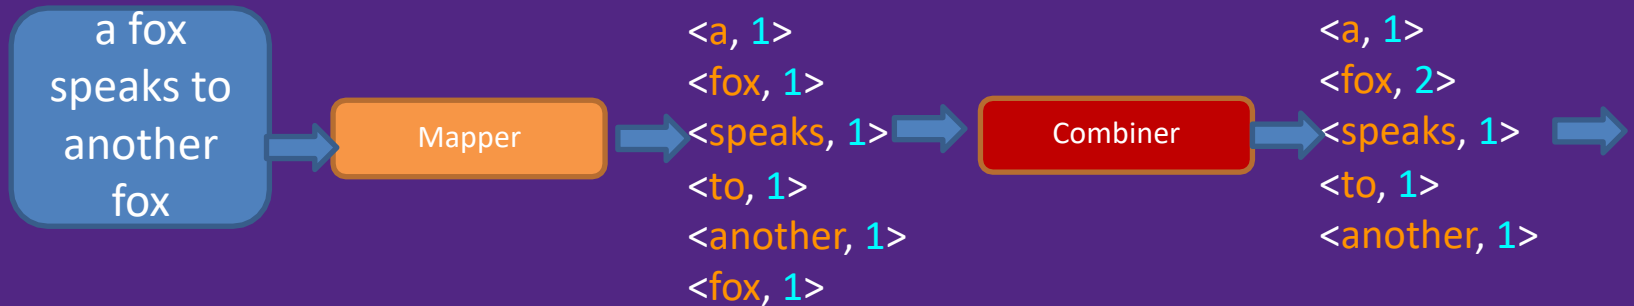
Use Combiners

- Perform local aggregation on the output of the map function but before the shuffle/sort phase
 - Word Count: Count occurrence of each word locally resulting in a pair (w, n) representing n occurrences of the word w
 - The number of intermediate pairs that go to the shuffle/sort is the number of unique words from that map
- This works when the operation on the pairs is associative and commutative

Use Combiners

- Hadoop Combiner is an optional class in the MapReduce framework
 - The combiner receives intermediate pairs from mappers
- You can save networking time by doing local aggregation
- When might this not be a good idea?

Combiner Example



- The reducer receives pairs (`w`, `n`) and sums all the `n`'s
- You can also incorporate the local aggregation in the map function

Word Count: Improved Mapper

Mapper (doc id a , doc d)

```
1       $H \leftarrow \text{New Associative Array (dictionary)}$ 
2      for all term  $t \in d$  do
3           $H\{t\} \leftarrow H\{t\} + 1$ 
4      for all term  $t \in H$  do
5          emit ( $t$ ,  $H(t)$ )
```

Make use of an associative array (or dictionary type) which is an abstract data type that is composed of a collection of (key,value) pairs such that a key only appears once

Hadoop Ecosystem

Other available tools

Other Tools

- Hive
 - SQL Database on Hadoop
- HBase
 - NoSQL Database on Hadoop
- Flume
 - Log processing and storage on Hadoop
- And more...

Summary

- MapReduce on Hadoop
 - Optimization – mappers, reducers, Combiners
- Other Hadoop Tools
 - Hive – SQL
 - HBase – NoSQL
 - Flume – Log Processing